



Mixup: Beyond Empirical Risk Minimization



New (Virtual) Labels



$$0.6 * \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$


$$+ 0.4 * \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$\Rightarrow \begin{bmatrix} 0.6 \\ 0.4 \end{bmatrix}$$



Multiclass:

$$0.6 * \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} + 0.4 * \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix} \Rightarrow \begin{bmatrix} 0.6 \\ 0 \\ 0 \\ 0.4 \\ 0 \end{bmatrix}$$



New hyperparameter (**alpha**)
alpha = 0.6
(1-alpha) = 0.4

Mixup

```
for (x1, y1), (x2, y2) in zip (loader1, loader2):
```

```
    lam = numpy.random.beta(alpha, alpha)
```

```
    mix = lam * x1 + (1. - lam) * x2
```

```
    y = lam * y1 + (1. - lam) * y2
```

```
    optimizer.zero_grad()
```

```
    output = net(mix)
```

```
    loss(output, y).backward()
```

```
    optimizer.step()
```





Mixup

- Better accuracy. (next slides)
- Increases the robustness to adversarial examples. (next slides)
- Reduces the memorization of corrupted labels.
- Stabilizes the training of generative adversarial networks.





Mixup: Classification

- If the dataset contains incorrect labels (labeling a cat as dog), mixup can help reducing the effect of incorrect labels.
- The model will not see a pure cat nor a pure dog during training!! (the training acc will be low, but we consider the val acc only).
- Don't mix images from the same class!






Results on CIFAR10 (200 epochs)

	Dataset	Model	ERM	Mixup (alpha=1)	Mixup (alpha=0.4)
paper	CIFAR10	ResNet18 (cifar10 variant)	94.4%	95.8%	-----
ours	CIFAR10	ResNet18 (cifar10 variant)	94.8%	95.9%	94.72%




train acc
tag: train acc



RED: ERM

val acc
tag: val acc



BLUE: Mixup



Mixup: Not only Images Data

- We can go beyond images and use mixup with tabular dataset in the same way.
- 3-layers MLP , IRIS dataset , 10 epochs

IRIS dataset	ERM	mixup
ours	83%	86%
paper	79%	83%



Robustness To Adversarial Examples

- Mixup can significantly improve the robustness of the neural network to adversarial examples..
- This gradient is important to find the adversarial examples (gradient ascend /next slide).
- (Hypothesis) It makes the norm gradient of the loss w.r.t the input small.
- **FAST GRADIENT SIGN METHOD:** computes the gradients of a loss function w.r.t the input image and then takes the direction that ascends the gradients to create a new image (not weights) that *maximizes* the loss.



Robustness To Adversarial Examples

Iterative Fast Gradient Sign Method: update the image not the weights

for n_steps:

for inputs, y in data:


input.requires_grad=True #very important


output = net(input)

loss(output, y).backward()

inputs = inputs + eps* inputs.grad.sign() #gradient ascend

- Test the model against the new data



	Attack	Model	eps	ERM	Mixup (alpha=1)
ours	White box attack	ResNet18	0.007 0.5 4	51% 20% 9%	46.% 16% 9%
ours	Black box attack	ResNet18	0.007 0.5 4	52% 36% 27%	46% 32% 24%
paper	White box attack	ResNet101	4	9.3%	14.8%
paper	Black box attack	ResNet101	4	43%	54%

Results After one
Iteration

Different results from the paper! We used ResNet18 not ResNet101! **Shallower model?**



Attack	Model	eps	ERM	Mixup (alpha=1)
White box attack	ResNet18	0.007	0.5%	1.25%
Black box attack	ResNet18	0.007	3.5%	8.4%

Results After 10
Iterations

Mixup Stabilizes the training of GANs

- Train the generator as usual.
- Mix a real image and a fake image and feed the new image to the discriminator

ERM:
$$\max_g \min_d \mathbb{E}_{x,z} \ell(d(x), 1) + \ell(d(g(z)), 0),$$

Mixup:
$$\max_g \min_{d, \lambda} \mathbb{E}_{x,z,\lambda} \ell(d(\lambda x + (1 - \lambda)g(z)), \lambda).$$



Conclusion

- **mixup** is a data augmentation method that consists of only two parts: 1-**random convex combination of raw inputs**, and correspondingly, 2-**convex combination of one-hot labels**.
- Another idea that gives worse results is to mix the **feature maps** instead of mixing the the raw inputs.
- Another idea also is to mix the **inputs only** (worse results).
- Using the combination of more than 2 images doesn't provide better results (2 images is enough).
- Not only images!!





The End





Questions 😄