



Mixup: Beyond Empirical Risk Minimization



New (Virtual) Labels



$$0.6 * \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$


$$+ 0.4 * \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$\Rightarrow \begin{bmatrix} 0.6 \\ 0.4 \end{bmatrix}$$



Multiclass:

$$0.6 * \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} + 0.4 * \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix} \Rightarrow \begin{bmatrix} 0.6 \\ 0 \\ 0 \\ 0.4 \\ 0 \end{bmatrix}$$



New hyperparameter (**alpha**)
alpha = 0.6
(1-alpha) = 0.4

Mixup

```
for (x1, y1), (x2, y2) in zip (loader1, loader2):
```

```
    lam = numpy.random.beta(alpha, alpha)
```

```
    mix = lam * x1 + (1. - lam) * x2
```

```
    y = lam * y1 + (1. - lam) * y2
```

```
    optimizer.zero_grad()
```

```
    output = net(mix)
```

```
    loss(output, y).backward()
```

```
    optimizer.step()
```





Mixup

- Better accuracy. (next slides)
- Increases the robustness to adversarial examples. (next slides)
- Reduces the memorization of corrupted labels.
- Stabilizes the training of generative adversarial networks.





Mixup: Classification

- If the dataset contains incorrect labels (labeling a cat as dog), mixup can help reducing the effect of incorrect labels.
- The model will not see a pure cat nor a pure dog during training!! (the training acc will be low, but we consider the val acc only).
- Don't mix images from the same class!






Results on CIFAR10 (200 epochs)

	Dataset	Model	ERM	mixup
paper	CIFAR10	ResNet18 (cifar10 variant)	94.4	95.8
ours	CIFAR10	ResNet18 (cifar10 variant)	94.8	95.9




train acc
tag: train acc



RED: ERM

val acc
tag: val acc



BLUE: Mixup

Robustness To Adversarial Examples

- Mixup can significantly improve the robustness of the neural network to adversarial examples..
- This gradient is important to find the adversarial examples (gradient ascend /next slide).
- (Hypothesis) It makes the norm gradient of the loss w.r.t the input small.
- **FAST GRADIENT SIGN METHOD:** computes the gradients of a loss function w.r.t the input image and then takes the direction that ascends the gradients to create a new image (not weights) that *maximizes* the loss.



Robustness To Adversarial Examples

Iterative Fast Gradient Sign Method: update the image not the weights

for n_steps:

for inputs, y in data:


input.requires_grad=True #very important


output = net(input)

loss(output, y).backward()

inputs = inputs + eps* inputs.grad.sign() #gradient ascend

- Test the model against the new data



	Attack	Model	eps	ERM	mixup
ours	White box attack	ResNet18	0.007 0.5 4	51% 20% 9%	46.% 16% 9%
ours	Black box attack	ResNet18	0.007 0.5 4	52% 36% 27%	46% 32% 24%
paper	White box attack	ResNet101	4	9.3%	14.8%
paper	Black box attack	ResNet101	4	43%	54%

Results After
one Iteration

Different results from the paper! We used ResNet18 not ResNet101! **Shallower model?**



Attack	Model	eps	ERM	mixup
White box attack	ResNet18	0.007	0.5%	1.25%
Black box attack	ResNet18	0.007	3.5%	%


Results After 10
Iterations

Mixup Stabilizes the training of GANs

- Train the generator as usual.
- Mix a real image and a fake image and feed the new image to the discriminator

ERM:
$$\max_g \min_d \mathbb{E}_{x,z} \ell(d(x), 1) + \ell(d(g(z)), 0),$$

Mixup:
$$\max_g \min_{d, \lambda} \mathbb{E}_{x,z,\lambda} \ell(d(\lambda x + (1 - \lambda)g(z)), \lambda).$$



Conclusion

- **mixup** is a data augmentation method that consists of only two parts: 1-**random convex combination of raw inputs**, and correspondingly, 2-**convex combination of one-hot labels**.
- Another idea that gives worse results is to mix the **feature maps** instead of mixing the the raw inputs.
- Another idea also is to mix the **inputs only** (worse results).
- Using the combination of more than 2 images doesn't provide better results (2 images is enough).
- Not only images!!





The End





Questions 😄